

# miniDB Documentation (v4 - 2022)

The miniDB project is a minimal and easy to develop rDBMS tool written in Python 3. MiniDB's main goal is to provide the user with as much functionality as possible while being easy to understand and even easier to expand. Thus, miniDB's primary market is students and researchers who aim to work with a tool that they can understand throughout, while being able to implement additional features as quickly as possible.

The miniDB project consists of several files, with the ones that encapsulate the main ideas and functionality being the following:

- `mdb.py`
- `miniDB/database.py`
- `miniDB/table.py`

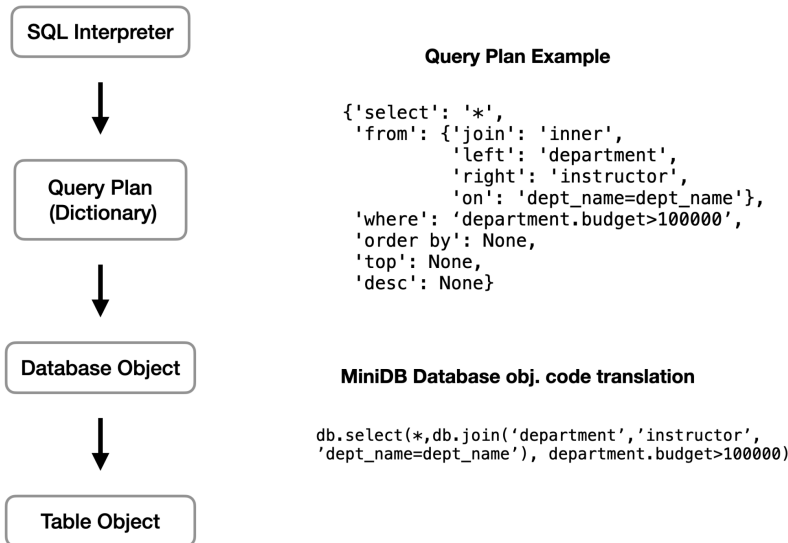
Other files include:

- `miniDB/btree.py`: BTree implementation.
- `miniDB/joins.py`: Implementations of Join algorithms.
- `miniDB/dashboard.py`: Dashboarding app that shows changes in tables live.
- `miniDB/misc.py`: Miscellaneous helper functions.

The structure of the documentation is as follows: Section 1 introduces the workflow, from the SQL interpretation to the access of data at the table level. Section 2 presents mSQL (miniDB's SQL) and a mapping of its functionality with respect to PostgreSQL. Section 3 provides more details about the source code and the main .py files

## 1. Introduction

The workflow that corresponds to the way data is stored and served using miniDB is presented in the following figure.



MiniDB's three main interfaces are:

1. The SQL interpreter that is implemented in the mdb.py file.
2. The database interface that is implemented in the miniDB/database.py file and acts as an intermediary between interfaces 1 and 3.
3. The table interface that handles all the per table functionality.

## 2. mSQL (miniDB SQL)

Like every RDBMS tool, miniDB supports the use of a slightly altered SQL implementation. This implementation (called SQL) has many features in common with PostgreSQL. The following table is a side by side comparison between PostgreSQL and mSQL, consisting of four parts: data description language (DDL) at database- and table-level, data manipulation language (DML), and locking functionality. As a comment, most of the table related SQL commands are identical, with slight changes regarding JOIN queries whereas database-related commands are different, with miniDB ones being closer to the psql commands.

Action	PostgreSQL	miniDB code (python3)
<b>DDL - Database level</b>		
Create DB	CREATE DATABASE name	cdb name # load if exists, create if not # saves are automatic
Save DB	SAVE DATABASE name	
Connect to DB	\c name	

Delete DB	DROP DATABASE name	rmdb name
List all DBs	\l	lsdb
List all tables in DB	\dt	lstb
<b>DDL - Table level</b>		
Create table	CREATE TABLE name (column_name1 column_type1 (primary key), ...)	CREATE TABLE name (column_name1 column_type1 (primary key), ...)
Delete table	DROP TABLE name	DROP TABLE name
Change column type	ALTER TABLE name ALTER COLUMN column TYPE INTEGER(int)	CAST column FROM name TO int
Import table from .csv file	CREATE TABLE name (column_name1 column_type1, ...); COPY name FROM file.csv;	IMPORT TABLE name FROM file.csv
Export table to .csv file	EXPORT TABLE name TO csvname.csv;	EXPORT TABLE name TO csvname.csv
Create index	CREATE INDEX name ON table(column) using BTREE;	CREATE INDEX name ON table(column) using BTREE;
Delete index	DROP INDEX name;	DROP INDEX name;
<b>DML</b>		
Insert record into table	INSERT INTO table VALUES (row);	INSERT INTO table VALUES (row);
Delete record(s) from table	DELETE FROM table WHERE condition;	DELETE FROM table WHERE condition;
Update values in table	UPDATE table SET column=value WHERE condition;	UPDATE table SET column=value WHERE condition;
Select-from-where over one table	SELECT [DISTINCT] columns/* FROM table WHERE condition [limit k] [ORDER BY column [ASC DESC]] [limit K]	SELECT [DISTINCT] columns/* FROM table WHERE condition [limit k] [ORDER BY column [ASC DESC]] [limit K]
Select-from-where over two tables	SELECT [DISTINCT] columns/* FROM table1 INNER JOIN table2 ON condition WHERE ...;	SELECT [DISTINCT] columns/* FROM table1 INNER JOIN table2 ON table1_column=table2_column

		WHERE ...;
Select-from-where over three (or more) tables	SELECT [DISTINCT] columns/* FROM table1 INNER JOIN table2 ON condition1 INNER JOIN table3 ON condition2 WHERE ...	SELECT [DISTINCT] columns/* FROM table1 INNER JOIN (SELECT [DISTINCT] columns/* FROM table2 INNER JOIN table3 ON table2_column=table3_column) ON table1_column=table[2/3].table[2/3] _column
<b>Locking</b>		
Lock/unlock table	BEGIN WORK; LOCK TABLE name IN EXCLUSIVE MODE; ... COMMIT WORK;	LOCK TABLE name; ... UNLOCK TABLE name;

(\*) Note about joins: (i) apart from "inner", mSQL also supports the following types of joins: "outer", "left", "right" ; (ii) three join algorithms have been implemented, namely "nlj" (nested loop join), which is the default, "smj" (sort merge join) and "inlj" (index nested loop join). In miniDB you can force the usage of a specific algorithm by using the keyword in front of the "join" part i.e. "sm join" or "inl join").

### 3. Codebase

#### mdb.py - mSQL Interpreter

The mdb.py file is the file that is run in order to create a connection with an existing miniDB database or to create a new one. This file is also the one that contains all the logic with regards to interpreting mSQL commands. To start a miniDB instance, run:

```
DB=DB_NAME python3.9 mdb.py
```

You need to specify a name for the database that you wish to connect to. If a database with that name does not exist, it will be created. The REPL window will look something like this:

```
(_)      (_)|__ \ |__ \
```

```
(DB_NAME) >
```

The mSQL interpreter works in a very specific and detached way with respect to the rest of the miniDB code. Its single job is to create a query plan in the form of a dictionary. In order to inspect the query plan of a specific query, the user should prefix the mSQL query with the EXPLAIN keyword. The result will look something like this:

```
{'select': '*',
 'from': {'join': 'inner',
          'left': 'instructor',
          'right': 'department',
          'on': 'dept_name=dept_name'},
 'where': None,
 'order by': None,
 'limit': None,
 'desc': None}
```

A very important concept that needs to be understood is the recursive nature of the Query Plan. As it is presented above, a query plan is a dictionary (a key-value store). Each key specifies an argument that will be supplied to the underlying python code in order for the mSQL command to be executed. However, if an argument is a new dictionary (like in the example above where the JOIN clause is a separate dictionary that is the value of the 'FROM' key), this nested dictionary is evaluated separately. When it produces a result (a Table object - will be discussed later on), the result takes the place of the nested dictionary. Then, the "parent" dictionary (in this case the 'SELECT' command) is evaluated and run, with its 'FROM' key having a Table object as value. This interim step can be thought of as a new dictionary that looks something like this:

```
{'select': '*',  
  'from': TABLE_OBJECT,  
  'where': None,  
  'order by': None,  
  'limit': None,  
  'desc': None}
```

This way, the interpreter can interpret infinitely nested commands, providing advanced functionality for a myriad of use-cases (multijoins, interim tables “view-like” tables etc.).

## miniDB/database.py - Database class

The Database class is the backbone of the miniDB pipeline. It is responsible for orchestrating all the commands that will be run in a specific database, while updating all the metadata that is needed to keep miniDB consistent and usable.

Each mSQL action (‘SELECT’, ‘INSERT INTO’ etc.) maps to a function of the Database class with the same name (if the action keyword is spaced - like ‘INSERT INTO’, the space is replaced with an underscore - mSQL ‘INSERT INTO’ maps to Database.insert\_into).

Managing the underlying metadata is the other important task that the Database class is responsible for. Currently, 4 meta\_tables are maintained:

- meta\_length -> keeps track of the number of rows of each table in the database (meta\_tables are excluded)
- meta\_locks -> keeps track of all the active locks in the database (meta\_tables are excluded)
- meta\_insert\_stack -> Keeps track of all the empty indexes in a table (result of a delete operation), in order to fill them in case a new row is inserted (if no indices then the new row is appended at the end of the table) (meta\_tables are excluded)
- meta\_indexes -> Keeps track of all the available indexes, as well as their name, the table and the column that they index, as well as their type (only B-Tree is currently supported).

## miniDB/table.py - Table class

The Table class is the class that is responsible for executing each operation that is specified by the Database class. It contains functions that do not follow strict naming rules, since they are not mapped in the same way that the Database functions are mapped to the mSQL interpreter. Each Table function executes a specific command in a specific way, meaning that extending the functionality of the miniDB project (when that involves creating new features) starts by creating a new Table function that implements the new feature. For example, unlike the unified DB.select

function, the Table class implements multiple select functions (Table.\_select\_where and Table.\_select\_where\_with\_btree). In this case, the Table.\_select\_where function implements a linear search for a given condition and the Table.\_select\_where\_with\_btree implements a B-Tree based search.